

LLVM Tutorial

Gregory J. Duck

Overview

- Basic LLVM tutorial
 - LLVM architecture overview (frontend, optimizer, linker)
 - LLVM *Intermediate Representation* (IR) overview
 - How go generate the IR
 - Programming with LLVM IR
- Basic LLVM knowledge is required for Assignment 2

LLVM Compiler System

- LLVM = ~~Low Level Virtual Machine~~ LLVM = LLVM
- The LLVM *Compiler Infrastructure Project*
 - Provides reusable components for building compilers
 - Reduce the time/cost to build a new compiler
 - Build static compilers, JITs, trace-based optimizers, ...
- The LLVM *Compiler Framework*
 - End-to-end compilers using the LLVM infrastructure
 - C and C++ are robust
 - Emit C code or native code for X86, Sparc, PowerPC, etc.

Three primary LLVM components

- The LLVM *Virtual Instruction Set* (Main focus of today's tutorial)
 - The common language- and target-independent *Intermediate Representation* (IR)
- A collection of *integrated libraries*
 - Analyses, optimizations, code generators, JIT compiler, profiling, etc.
- A collection of *tools* built from the libraries
 - Assemblers, debugger, linker, code generator, compiler driver, optimizer, etc.

The LLVM C/C++ Compiler

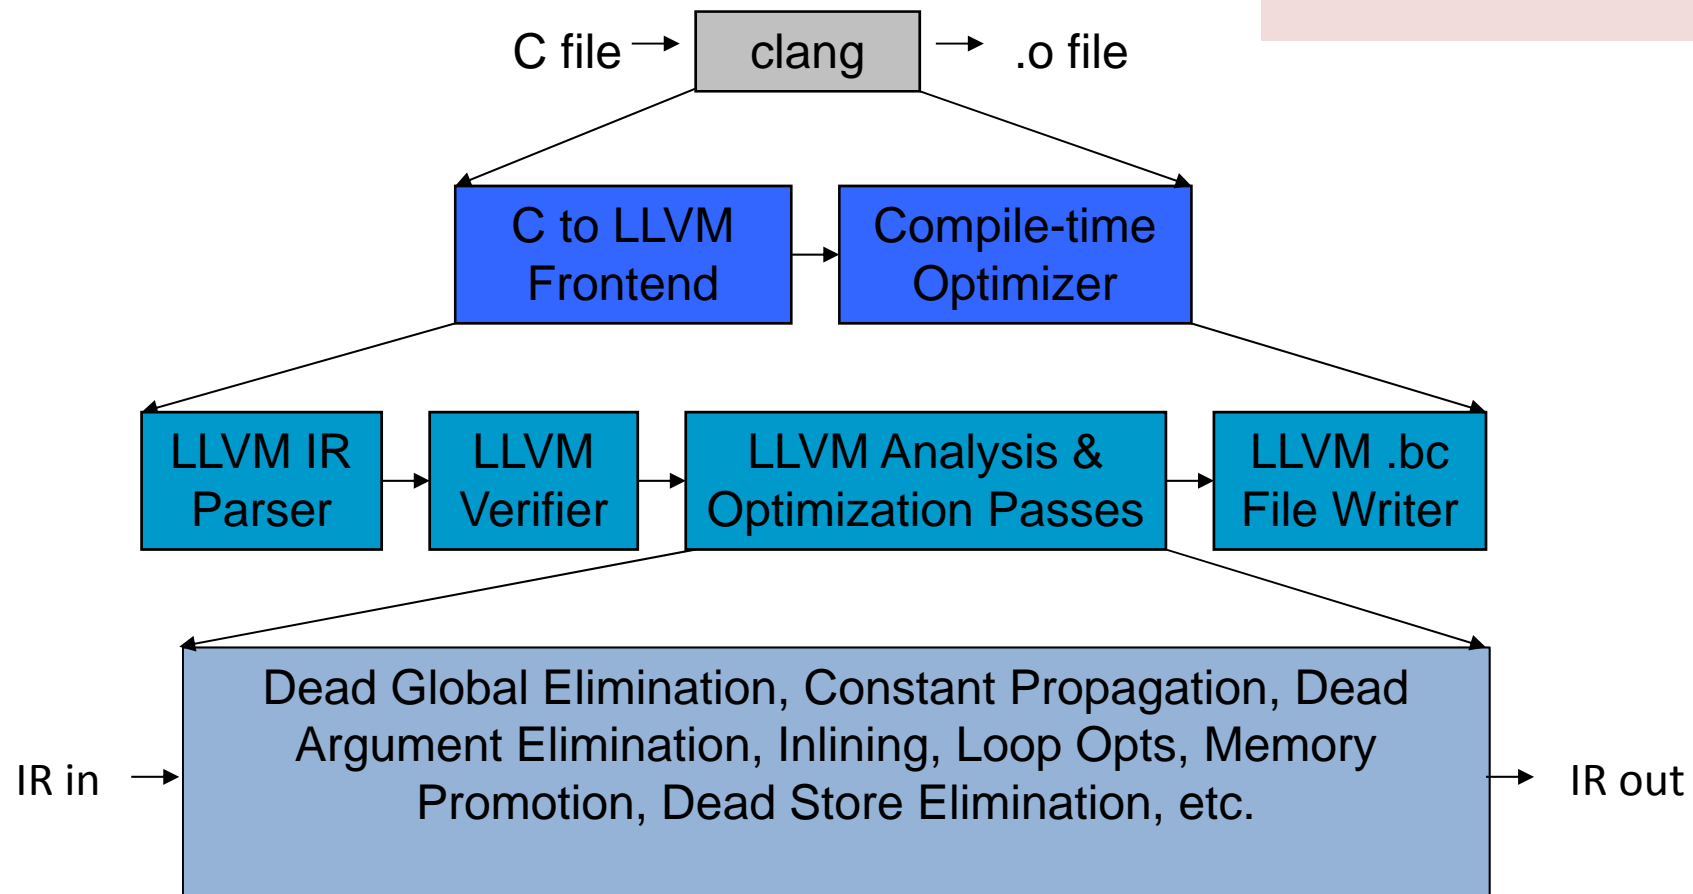
- LLVM C/C++ is a standard compiler:
 - Compatible with GCC options and extensions
 - Makefile compatible

```
$ clang -c hello.c  
$ clang -o hello hello.o  
$ ./hello  
Hello World
```

- Standard compiler organization, which uses LLVM as midlevel IR:
 - Language specific front-end lowers code to LLVM IR
 - Language/target independent optimizers improve code
 - Target specific back-end (e.g. x86) generate native code

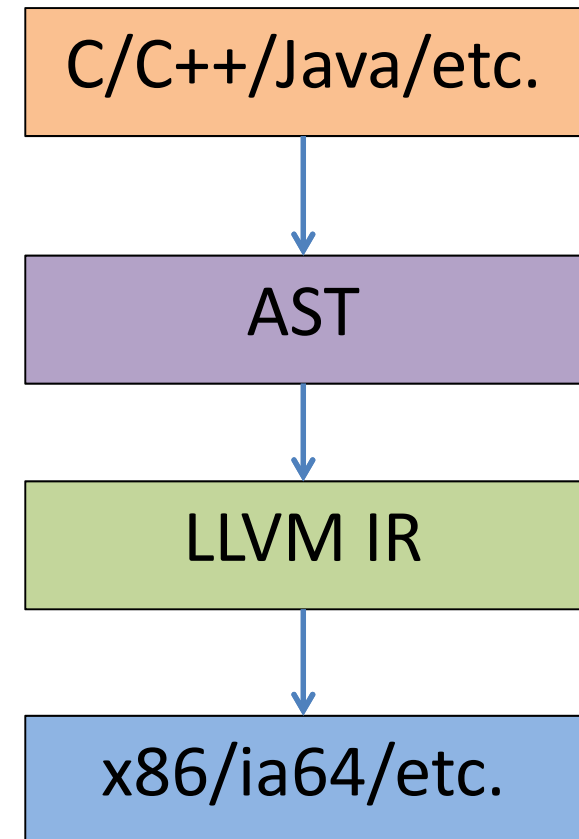
LLVM Compilation Overview

```
clang -c hello.c
```




Goals of LLVM IR

- Language- and Target-Independent
 - AST-level IR is not language independent
- One IR for analysis **and** optimization
 - IR must be able to support aggressive IPO, loop opts, scalar opts, ... high- **and** low-level optimization!
- Optimize early as possible
 - No lowering in the IR!



LLVM IR Program Structure

- Module contains Functions/GVs
 - Module is unit of compilation
- Function contains BasicBlocks
 - Functions correspond to C functions
- BasicBlock contains instructions
 - Each block ends in a control flow instruction
- Instruction is opcode + operands



My slides go “bottom up”

DEMO: Generating LLVM IR

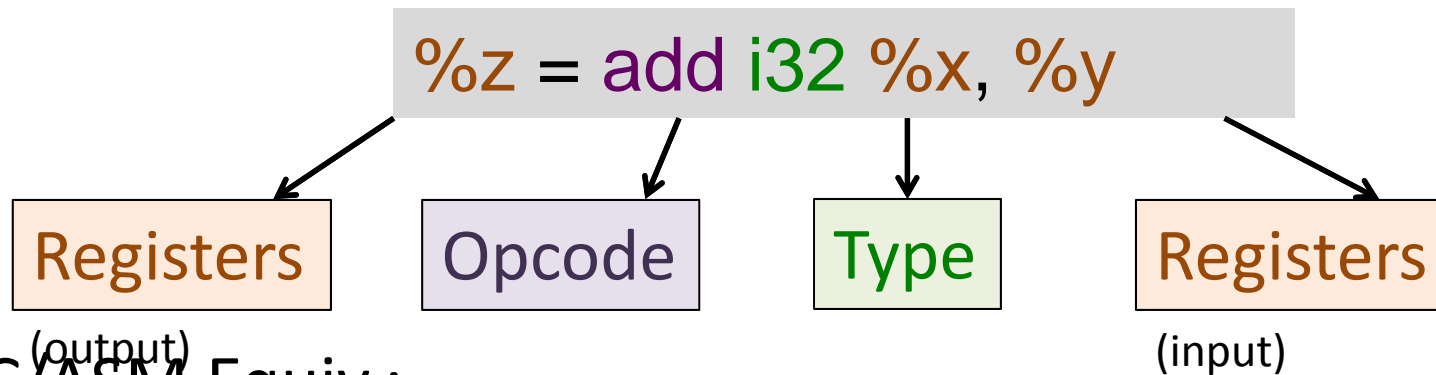
LLVM IR Instruction Basics

- LLVM IR Instructions are of the general form:

`retval = opcode type operand1, . . . , operandn`

(There can be other forms!)

- For example, addition:



- C/ASM Equiv.:

```
int x, y, z; . . . ; z = x + y;
```

```
add %rbx, %rax
```

LLVM IR Instruction Components

```
%z = add i32 %x, %y
```

- **Registers:**
 - Register names begin with a `%`, e.g. `%x`, `%y`, `%z`, etc.
 - Registers are *unlimited* (unlike real machines, e.g. x86).
- **Types:**
 - Integer types, e.g. `i32` = 32bit integer.
 - Other types: `float` (floating point), `i32*` (pointers), `%struct.node` (structures), `[100 x i32]` (arrays), etc.

LLVM IR Instruction Components

```
%z = add i32 %x, %y
```

- Opcodes (various):
 - *Arithmetic*: add, sub, mul, div, etc. *Bitwise*: lshr, asll, and, or, xor, etc. *Comparison*: cmp *Control*: ret, call, br, switch, etc. *Casts*: bitcast, zext, sext, trunc, ptrtoint, inttoptr, etc. *Memory*: alloca, load, store, getelementptr
- Return value is sometimes optional.

```
ret i32 %x
```

```
call @func(i32 %x, i32 %y)
```

- Some opcodes have special syntax.

```
%z = bitcast i32* %y to i16*
```

LLVM IR Instruction Components

- Instruction operands can also be *constants*:

```
%z = add i32 %x, 1
```

$z = x + 1$

- Also *global variables*:

$z = (\text{short } *)\text{inc} ; // \text{Global}$

```
%z = bitcast i32* @inc to i16*
```

- Local (registers) start with a %
- Globals start with a @, e.g. @inc
- Note that functions are typically globals, e.g. @main

LLVM IR Key Properties

- Low-level and target-independent semantics
 - RISC-like three address code / SSA form
 - Infinite virtual register set
 - Strongly typed
 - Simple, low-level control flow constructs
 - Load/store instructions with typed-pointer

Properties: LLVM IR is in SSA Form

- LLVM is in *Single Static Assignment* form:
 - Values are created once and can never change.
 - Motivation: simplifies analysis and optimization.
 - LLVM has infinite registers, so not a problem.
 - E.g. this instruction is **not** in SSA form (%x redefined):

```
%x = add i32 %x, %y  
ret i32 %x
```

← Error

```
%1 = add i32 %x, %y  
ret i32 %1
```

OK!

(SSA is standard practice)

(Register names can be numbers, e.g. %1, %2, etc.)

Properties: LLVM IR is Strongly Typed

- LLVM is *strongly typed*:

```
%z = add i32 %x, %y  
%w = add i64 %z, %z
```

← Type error

(%z is int32 and passed to int64 instruction)

- Types must be explicitly converted via instructions:

```
%y = zext i32 %x to i64
```

int32 to int64 (zero extend)

```
%y = trunc i64 %x to i32
```

int64 to int32 (truncate)

```
%y = bitcast i32* %x to i64*
```

int32* to int64*

(others for int2ptr, ptr2int, int2float, etc., etc.)

LLVM Instruction Types

- LLVM has many instruction types:
 - Arithmetic & Bitwise
 - Casts
 - Memory
 - Function calls
 - Pointer arithmetic
 - Comparison and Control Flow
 - PHI-nodes
- All C/C++ compile to these basic instructions.

LLVM Instruction Types: Memory

- Memory is accessed via two instructions:

```
%x = load i32 %y
```

$x = *y$

```
store i32 %y, i32* %x
```

$*x = y$

- All other instructions must work on registers (unlike x86)

- Stack memory allocated via *alloca*

```
%x = alloca [100xi32]
```

`int x[100]`

LLVM Instruction Types: GEP

- GEP = *Get Element Pointer*
- GEP is for *pointer arithmetic*

$y = x + 10$

$y = x + 10$

```
%y = getelementptr i32* %x, i32 10
```

$y = \&x \rightarrow \text{next}$ a.k.a. $y = \&x[0].\text{next}$

$y = x + 0 + \text{offsetof}(\text{next})$

```
%y = getelementptr %struct.node* %x, 0, 1
```

- <http://llvm.org/docs/GetElementPtr.html>
- GEP = pointer arithmetic, nothing else.

LLVM Basic Blocks & Control Flow

- Instructions are arranged into *Basic Blocks*:

```
return x * y + z;
```



```
label:  
%1 = mul i32 %x, %y  
%2 = add i32 %1, %z  
ret i32 %2
```

- Each Basic Block:
 - Has a label (e.g. %label)
 - A sequence of *branch free* instructions.
 - Ends with a *terminator* instruction (e.g. ret)
 - (LLVM) Has exactly *one* entry and *one* exit (no “fall through” e.g. x86).
- Together forms the *Control Flow Graph* (CFG)

LLVM Instruction Types: Control Flow

- LLVM *branch instructions* connect Basic Blocks.

```
br i1 %x, label %true, label %false
```

Condition Var (Boolean)

True branch

False branch

```
br label %label
```

(unconditional branch)

- Return instruction: `ret i32 %x`
- Others: `invoke`, `unreachable`

LLVM Instruction Types: Comparison

- Conditional branches take a Boolean input:

```
br i1 %b, label %true, label %false
```

- Typically the results of a *comparison instruction*:

```
%b = icmp ult i32 %x, %y  
br i1 %b, label %true, label %false
```

- Multiple comparison operations supported: *eq, ne, gt, ge, lt, le* (both *signed* and *unsigned*)

LLVM Control Flow Example 1

```
entry:  
%1 = add i32 %x, 1  
%2 = cmp lt i32 %x, %y  
br i1 %2 label %T, label %F
```

```
if (x + 1 < y)  
    return x * y + z;  
return 0;
```

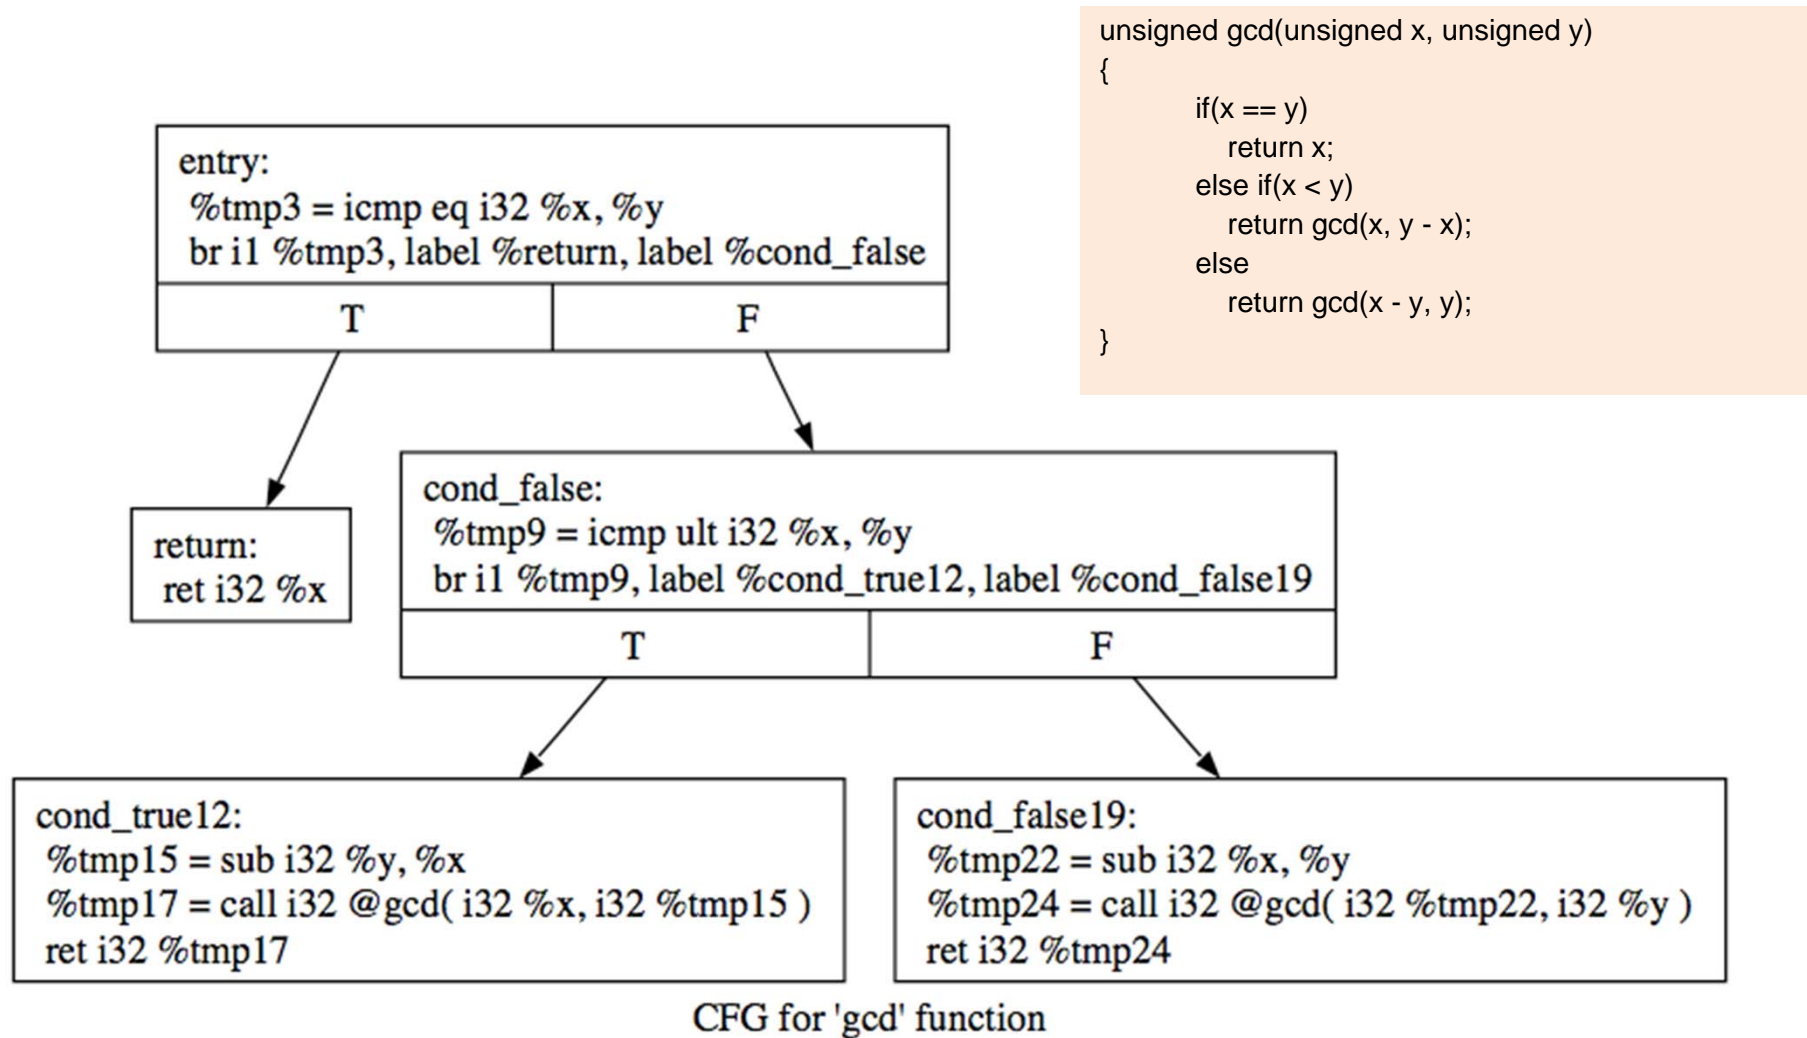
true

false

```
T:  
%3 = mul i32 %x, %y  
%4 = add i32 %3, %z  
ret i32 %2
```

```
F:  
ret i32 0
```

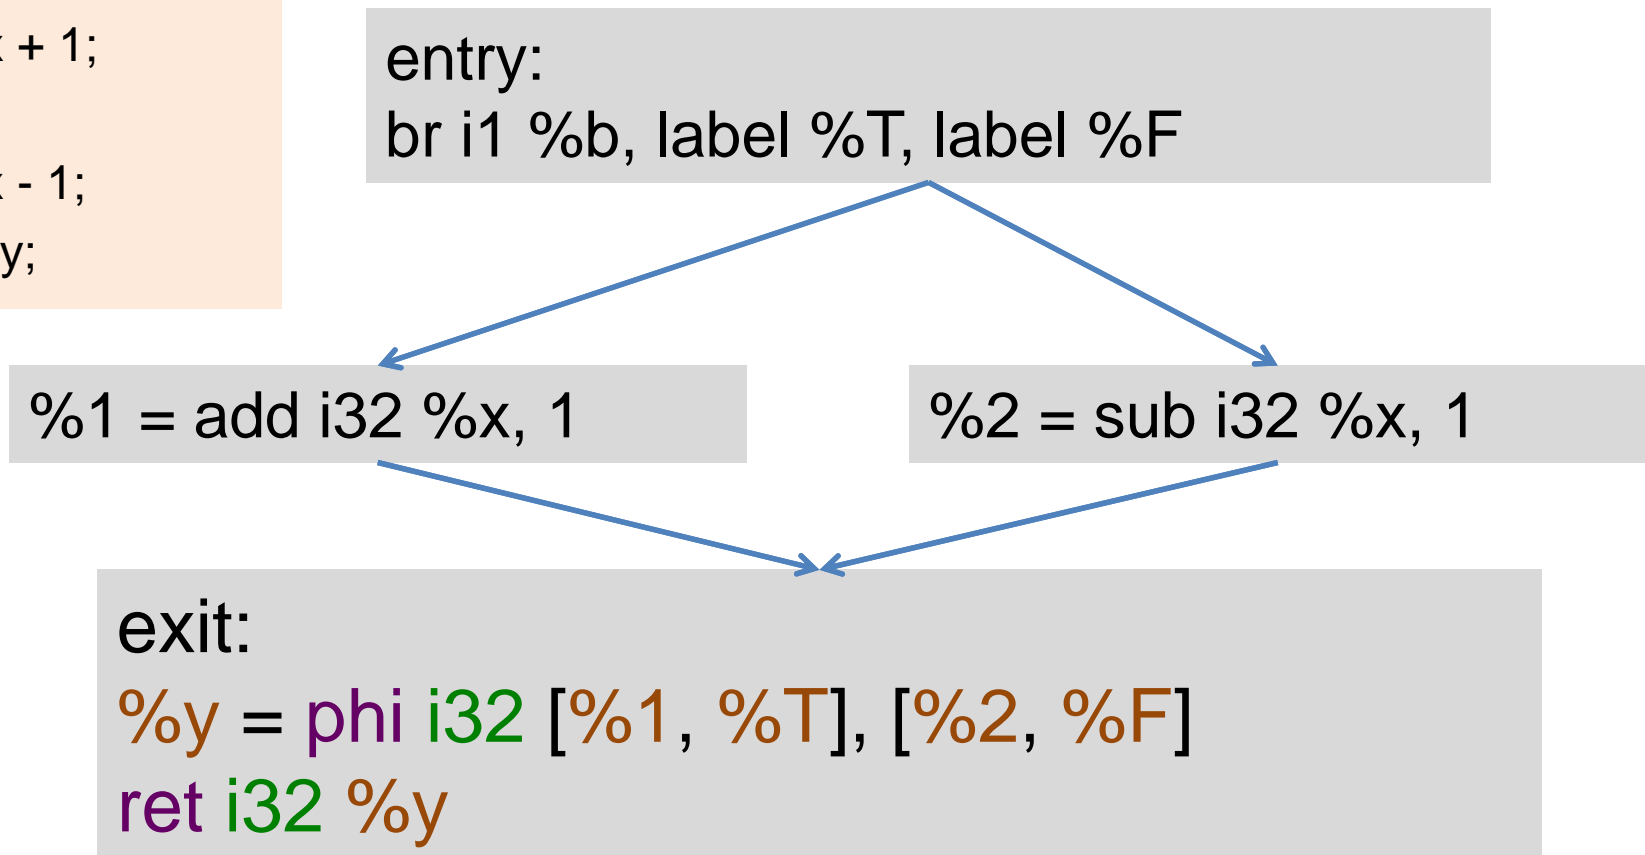
LLVM Control Flow Example 2



LLVM IR Instructions: PHI Nodes

- PHI Nodes are special instructions that “join” values:

```
if (b)
  y = x + 1;
else
  y = x - 1;
return y;
```



LLVM Functions

- LLVM functions roughly correspond to C functions:

```
int mul_add(int x,  
            int y,  
            int z)  
{  
    return x * y + z;  
}
```

```
define i32 @mul_add(i32 %x,  
                   i32 %y,  
                   i32 %z)  
{  
    entry:  
    %1 = mul i32 %x, %y  
    %2 = add i32 %1, %z  
    ret i32 %2  
}
```

- Each LLVM function has exactly one entry point

Part 2: Programming with LLVM IR

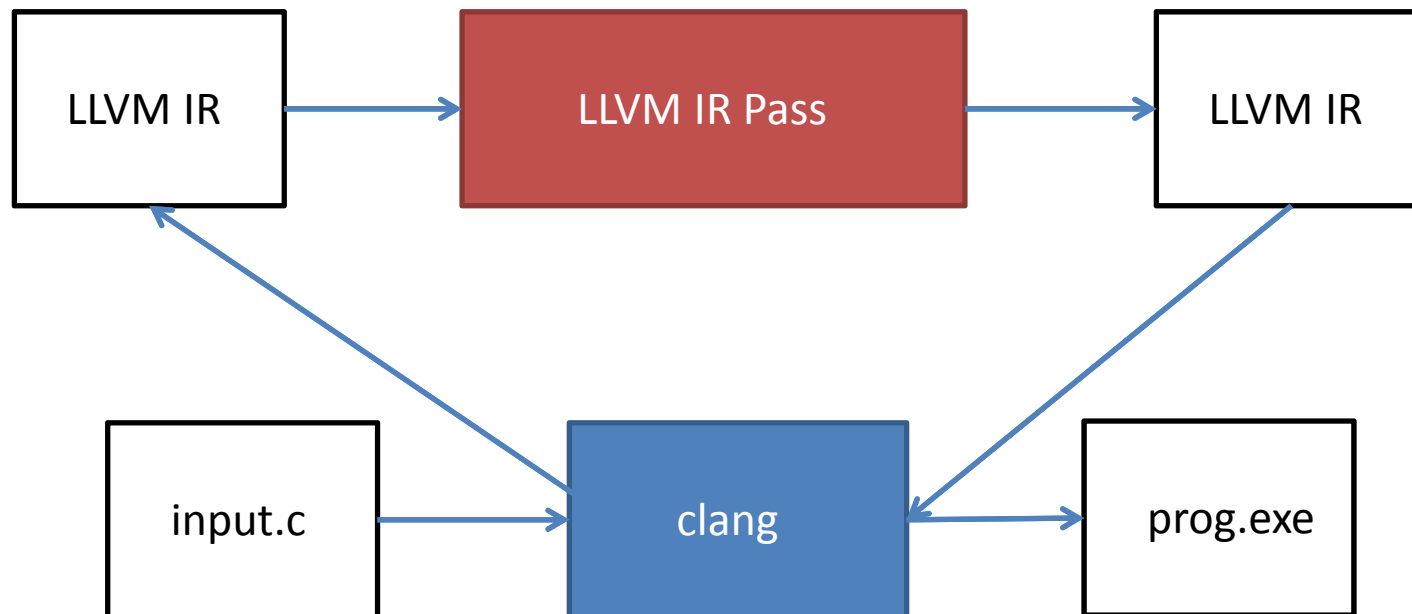
- LLVM is implemented in C++. (optional for assignment 2)
- To install:

```
sudo apt-get install llvm-dev
```

- Installs the LLVM API (or the IR) header files and libraries.
- The LLVM API can be used to:
 - Create new LLVM compiler passes (e.g. optimizations, analysis).
 - Create standalone programs that manipulate LLVM IR.

Option 1: Compiler Pass

- A compiler pass transforms LLVM IR:



- Con: lots of boilerplate to implement.

Option 2: Standalone C++

- An alternative to create a standalone C++ program:

```
#include "llvm/IRReader/IRReader.h"
...
int main(int argc, char **argv)
{
    // Step (1) Parse the given IR File
    LLVMContext &Context = getGlobalContext();
    SMDiagnostic Err;
    Module *M = ParseIRFile(argv[1], Err, Context);
    if (M == nullptr)
    {
        fprintf(stderr, "failed to read IR file %s\n", argv[1]);
        return 1;
    }
    ...
}
```

Standalone C++ (cont.)

- Reading in an IR file generates a *Module* object.
- Instructions can be traversed using C++ range loops:

```
int main(int argc, char **argv)
{
    ...
    // Step (2) Traverse all instructions
    for (auto &F: *M)           // For each function F
        for (auto &BB: F)       // For each basic block BB
            for (auto &I: BB)    // For each instruction I
                I.dump();        // Dump the instruction!
    ...
}
```

Compiling the Standalone

- Use the magic command to compile:

```
$ clang++-3.4 -o Test Test.cpp `llvm-config-3.4 --cxxflags`  
`llvm-config-3.4 --ldflags` `llvm-config-3.4 --libs` -lpthread  
-lncurses -ldl
```

- Run the command to dump the instructions:

```
$ Test mul_add.ll  
%1 = alloca i32, align 4  
%2 = alloca i32, align 4  
%3 = alloca i32, align 4  
store i32 %x, i32* %1, align 4  
store i32 %y, i32* %2, align 4  
...
```

Making it more Interesting

- The LLVM API makes heavy use of dynamic casting.
- E.g. to dump only call instruction names:

```
int main(int argc, char **argv)
{
    ...
    // Step (2) Traverse all instructions
    for (auto &F: *M)                // For each function F
        for (auto &BB: F)            // For each basic block BB
            for (auto &I: BB) // For each instruction I
            {
                CallInst *Call = dyn_cast<CallInst>(&I);
                if (Call == nullptr) continue;
                Function *G = Call->getCalledFunction();
                if (G == nullptr) continue;
                printf("Name = %s\n", G->getName().str().c_str());
            }
}
```


Standalone Version 2

- Use the magic command to re-compile:

```
$ clang++-3.4 -o Test Test.cpp `llvm-config-3.4 --cxxflags`  
`llvm-config-3.4 --ldflags` `llvm-config-3.4 --libs` -lpthread  
-lncurses -ldl
```

- Run the command to dump the instructions:

```
$ Test sha256.ll  
llvm.memcpy.p0i8.p0i8.i64  
sha256_init  
sha256_update  
sha256_final  
printf  
...
```

Questions?

Appendix: 2 Samples

```
#include <stdio>
#include "llvm/IR/LLVMContext.h"
#include "llvm/IR/Module.h"
#include "llvm/IR/Function.h"
#include "llvm/IR/Instructions.h"
#include "llvm/IR/Instruction.h"
#include "llvm/IRReader/IRReader.h"
#include "llvm/Support/SourceMgr.h"
using namespace llvm;
int main(int argc, char **argv)
{
    // Read the IR file.
    LLVMContext &Context = getGlobalContext();
    SMDiagnostic Err;
    Module *M = ParseIRFile(argv[1], Err, Context);
    if (M == nullptr)
    {
        fprintf(stderr, "error: failed to load LLVM IR file \"%s\\\"", argv[1]);
        return EXIT_FAILURE;
    }

    // Dump all instructions.
    for (auto &F: *M)
        for (auto &BB: F)
            for (auto &I: BB)
                I.dump();
    return 0;
}
```

```
... // add include files

int main(int argc, char **argv)
{
    // Read the IR file.
    LLVMContext &Context = getGlobalContext();
    SMDiagnostic Err;
    Module *M = ParseIRFile(argv[1], Err, Context);
    if (M == nullptr)
    {
        fprintf(stderr, "error: failed to load LLVM IR file \"%s\\\"", argv[1]);
        return EXIT_FAILURE;
    }
    // Dump all instructions.
    for (auto &F: *M)
        for (auto &BB: F)
            for (auto &I: BB)
            {
                CallInst *Call = dyn_cast<CallInst>(&I);
                if (Call == nullptr)
                    continue;
                Function *G = Call->getCalledFunction();
                if (G == nullptr)
                    continue;
                printf("%s\\n", G->getName().str().c_str());
            }
}
```